

Critic-Filtered Generative Data Augmentation for Low-Resource Chinese Character Classification: A C-WGAN-GP Study

Joseph Catanzarite

Johns Hopkins University - Whiting School of Engineering

EN.705.603 - Introduction to Generative AI

Spring 2026 - Final Submission

Abstract

Deep learning models for Chinese character recognition face a fundamental challenge: thousands of visually intricate character classes, each requiring labeled examples that are expensive to collect. This paper investigates generative data augmentation — training a generative model on scarce real data and sampling synthetic training examples from it — using the Chinese-MNIST dataset — a 15-class, 64×64 grayscale analogue of the classic MNIST (Modified National Institute of Standards and Technology) handwritten-digit benchmark — as a controlled proxy for low-resource domains such as medical imaging. The study trains a Conditional Wasserstein Generative Adversarial Network with Gradient Penalty (C-WGAN-GP) and evaluates its synthetic output as augmentation for a Convolutional Neural Network (CNN) classifier in two experiments. Experiment 1 (50 real images/class, unfiltered synthetic data) finds that augmentation *degrades* test accuracy, significantly so at a 4:1 synthetic-to-real ratio (-2.22% , $p = 0.008$). Experiment 2 (25 real images/class) applies this paper's original contribution — a quality filter that scores a large candidate pool with the trained critic and retains only the top 25% — and finds that filtered augmentation *significantly improves* accuracy, monotonically with ratio ($+5.55\%$ at 1:1, $p = 0.0085$; $+8.42\%$ at 4:1, $p = 0.0005$), recovering nearly half the accuracy lost to the deeper scarcity. All comparisons use paired t-tests across five seeds on identical held-out real test data. The central finding: generative augmentation is not plug-and-play — the same generator is harmful or strongly beneficial depending on whether its output is curated — and the adversarially-trained critic, normally discarded after training, provides an effective zero-cost quality gate. The three-architecture comparison proposed at midpoint — spanning variational, adversarial, and hybrid conditional generators — was descoped to this deep single-architecture study under compute and schedule constraints; the comparative study remains the natural extension.

1. Introduction

1.1 The Problem: Data Scarcity in a High-Cardinality Domain

The Chinese writing system is, by any measure, one of the most complex visual languages humans have ever produced. The Unicode CJK (Chinese, Japanese, Korean) block alone contains over 20,000 characters, each a unique configuration of strokes layered and balanced according to compositional rules developed over four millennia. For a machine learning model — specifically a Convolutional Neural Network (CNN), a deep learning architecture that learns to recognize patterns by scanning images through layers of learnable filters — recognizing all these characters reliably requires a large number of labeled training examples per class.

This is the core problem. Collecting and labeling thousands of handwritten examples per character is expensive and slow. When training data is sparse, CNNs tend to *overfit*: rather than learning the underlying geometric logic of a character — its stroke structure, its proportions, its radical components — the model memorizes the specific pixel patterns it has seen and fails to generalize to new examples. The result is a classifier that performs well on its training set but poorly in the real world.

One powerful remedy is **data augmentation**: artificially expanding the training set by creating new examples. Simple augmentation applies geometric transforms to existing images — rotating, flipping, cropping, adding noise. These are useful, but they are bounded by the original data: every augmented image is still a transformation of something that already existed. **Generative data augmentation** is more ambitious. Instead of transforming existing images, we train a generative model — a neural network that learns the underlying distribution of the data — and sample entirely new examples from that learned distribution. Done well, the synthetic images are not copies or distortions of training examples; they are novel, plausible instances that the model has, in a meaningful sense, imagined.

1.2 Three Architectures, One Question

Three families of conditional generative models have emerged as strong candidates for this task, each with a distinct philosophy:

The Conditional Variational Autoencoder (C-VAE) [1, 2] takes a probabilistic approach. It learns to compress images into a structured, low-dimensional *latent space* — a kind of internal coordinate system where similar characters cluster together — and then reconstructs images from points in that space. The “conditional” part means the model is always told which character class it is working with, so generation is class-specific. The

VAE is mathematically principled and trains stably, but its outputs can be slightly blurry because it optimizes a pixel-level reconstruction objective.

The Conditional Generative Adversarial Network (C-GAN) [3, 4] takes an adversarial approach. Two networks — a Generator that creates fake images and a Discriminator that tries to tell fakes from real ones — compete in a minimax game. Over time, the Generator learns to produce images realistic enough to fool the Discriminator. GANs tend to produce sharper, more photorealistic outputs than VAEs, but they are notoriously harder to train: the adversarial balance can tip, leading to *mode collapse* (where the Generator finds a few images that fool the Discriminator and stops exploring) or outright training instability.

The Conditional VAEGAN (C-VAEGAN) [5, 6] is a hybrid that aims to capture the best of both. It uses the VAE’s encoder-decoder structure to maintain a coherent latent space, but replaces the VAE’s blurry pixel-level reconstruction loss with the GAN’s perceptual “does this look real?” criterion. Theoretically, this should yield outputs that are both structurally coherent *and* visually sharp. Whether this advantage holds empirically, on Chinese character data specifically, is one of the central questions this paper investigates.

The midpoint proposal framed all three architectures for implementation on a shared dataset, evaluation on a shared metric suite, and comparison of downstream augmentation utility, under the deliberately open question: *which generative paradigm best serves the augmentation use case for Chinese character recognition?* The final study, as described in Section 1.4, executes the C-GAN arm of that design in depth — realized as the C-WGAN-GP — and the cross-architecture comparison remains the design’s natural completion.

1.3 Research Hypotheses

H₁ (Generation Quality): The C-VAEGAN will achieve a lower Fréchet Inception Distance (FID), a higher Structural Similarity Index Measure (SSIM), and better Learned Perceptual Image Patch Similarity (LPIPS) scores than either the C-VAE or C-GAN individually.

H₂ (Augmentation Efficacy): A CNN classifier trained on real data augmented with C-VAEGAN synthetic samples will achieve significantly higher validation accuracy than one trained on real data alone, with $p < 0.05$ on a paired t-test across five independent experimental runs.

H₃ (Quality Filter Value): The novel two-stage quality filter introduced in this paper will improve augmentation efficacy relative to using unfiltered synthetic samples, for all three generative architectures. (As Section 1.4 details, the final study tests H₃ for the C-GAN architecture; the cross-architecture portion is future work.)

1.4 Scope of the Final Study

The midpoint proposal specified a three-architecture comparison (C-VAE, C-GAN, C-VAEGAN). Under compute and schedule constraints, the final study descopes to a deep evaluation of one architecture — the C-GAN trained with the Wasserstein Gradient Penalty objective (C-WGAN-GP) — across two experiments that together isolate the role of synthetic-data quality. This descoping preserves the project’s most novel element: the quality filter proposed in Section 5.4 is implemented (Stage 1) and empirically validated, directly testing hypothesis H_3 for the C-GAN architecture. Hypothesis H_1 , which requires the cross-architecture comparison, is explicitly not tested and is carried forward as future work. Hypothesis H_2 is tested for the C-GAN in both filtered and unfiltered regimes, with results that turn out to depend critically on filtering — itself the paper’s central finding.

2. Related Work

2.1 Variational Autoencoders

The Variational Autoencoder (VAE), introduced by Kingma and Welling in 2014 [1], was a landmark contribution to generative modeling. At its core, a VAE is an encoder-decoder architecture: the encoder maps an input image to a probability distribution in a low-dimensional latent space, and the decoder reconstructs the image from a sample drawn from that distribution. The key insight is that by forcing the latent distribution to remain close to a standard normal distribution, the model learns a smooth, continuous latent space in which nearby points decode to visually similar images — a structure that makes controlled, interpolated generation possible.

Sohn, Lee, and Yan [2] extended this framework with conditioning, producing the Conditional VAE (C-VAE). By feeding the class label into both the encoder and decoder, the C-VAE can generate examples of a specific class on demand — exactly what is needed for a labeled augmentation pipeline.

2.2 Generative Adversarial Networks

Goodfellow and colleagues introduced Generative Adversarial Networks (GANs) in 2014 [3] with a deceptively simple idea: train two networks simultaneously. A Generator network G learns to map random noise vectors to realistic-looking images; a Discriminator network D learns to distinguish real images from G ’s fakes. Each improves by trying to beat the other. When training converges, G has learned to produce images indistinguishable from the real data distribution. Mirza and Osindero [4] introduced the conditional extension (C-GAN) by providing the class label as an additional input to both G and D , enabling class-specific generation.

Arjovsky and colleagues [13] later reframed adversarial training around the Wasserstein distance (the Wasserstein GAN, or WGAN), and Gulrajani and colleagues [7] stabilized it further with the Gradient Penalty technique (WGAN-GP), which dramatically improves GAN training by penalizing the Discriminator for violating a smoothness constraint — addressing one of the most common failure modes in adversarial training.

2.3 VAE-GAN Hybrids

Larsen and colleagues [5] proposed combining the two paradigms: use the VAE’s encoder-decoder as the generative backbone, but replace the VAE’s pixel-level reconstruction loss with a learned perceptual loss computed from the GAN Discriminator’s internal feature representations. The intuition is that pixel-level loss penalizes the model for every slightly misplaced pixel — leading to blurriness as a conservative hedge — whereas a perceptual loss asks only whether the image looks right overall. The conditional variant (C-VAEGAN) is directly applicable to Chinese character synthesis, where we need both structural regularity (favoring the VAE’s latent structure) and visual sharpness (favoring the GAN’s perceptual criterion).

2.4 Chinese Character Recognition and Synthesis

Deep CNN-based recognition of Chinese handwriting has advanced substantially, with state-of-the-art systems achieving over 95% accuracy on benchmark datasets [8]. However, these results depend on very large labeled datasets — the CASIA-HWDB (Center for Analysis and Statistics of Handwriting - Handwritten Database) corpus, for instance, contains nearly 3.9 million character images across 7,356 classes [9]. The Chinese-MNIST dataset [10], a more compact benchmark with 15 character classes and approximately 1,000 images each, provides a tractable starting point for generative augmentation experiments.

Kong and Xu [6] directly addressed Chinese character synthesis using a C-VAEGAN architecture, demonstrating feasibility on a small 200-samples-per-class regime. Their work encountered training instability — discriminator loss collapsing to near zero, and KL (Kullback-Leibler) divergence numerical instability after approximately eight training epochs — and reported no downstream augmentation evaluation. This paper extends their work in three directions: a systematic three-model comparison, a downstream classification evaluation, and the introduction of a quality filtering pipeline.

2.5 Positioning of This Work

To the author’s knowledge, no prior published work has conducted a controlled three-way comparison of C-VAE, C-GAN, and C-VAEGAN specifically for Chinese character data augmentation with downstream OCR (Optical Character Recognition) accuracy as the evaluation criterion. This

comparison directly addresses a gap in the literature: not which model generates the prettiest images, but which generative paradigm most usefully serves a downstream classification task.

3. Research Problem Statement

We begin with a labeled training set $D_{\text{real}} = \{(x_i, y_i)\}$ of Chinese character images, where each x_i is a grayscale image and each y_i is one of C character class labels. Our generative model — in the executed study, the C-WGAN-GP — is trained on this real data and used to produce a synthetic dataset D_{synth} . We then train a CNN classifier on the augmented set $D_{\text{real}} \cup D_{\text{synth}}$ and compare its performance to a baseline CNN trained on D_{real} alone. The central question is:

As proposed at midpoint: which conditional generative architecture — C-VAE, C-GAN, or C-VAEGAN — produces synthetic data that most effectively augments the real training set for downstream Chinese character classification? As executed in the final study: does C-WGAN-GP synthetic data augment that training set effectively, and does critic-based quality filtering change the answer?

This decomposes into three measurable sub-questions: (1) Which architecture produces the highest-quality synthetic characters, as measured by FID, SSIM, and LPIPS? (2) Which produces the greatest lift in CNN classifier accuracy when added to the real training data? (3) Does the novel two-stage quality filter improve augmentation efficacy, and does its benefit differ across architectures?

Primary dataset: Chinese-MNIST — 15 numeral-character classes: the ten digits zero through nine (零、一、二、三、四、五、六、七、八、九) plus the five magnitude characters for ten, hundred, thousand, ten thousand, and one hundred million (十、百、千、万、亿); 1,000 grayscale 64×64-pixel images per class, 15,000 images total.

Performance criteria: Generation quality (FID, SSIM, LPIPS), downstream CNN classification accuracy, and paired t-test statistical significance ($\alpha = 0.05$) across five independent runs.

4. Evaluation Metrics

Before describing the architectures, it is worth pausing to explain precisely how we will measure success. A generative model can fail in many ways: its images might look nothing like real characters, they might all look like the same character regardless of the conditioning label, or they might be

technically realistic but somehow not useful for training a classifier. Each of our three metrics captures a different dimension of this question.

4.1 Fréchet Inception Distance (FID)

The Fréchet Inception Distance, or FID, asks: how similar is the *distribution* of generated images to the distribution of real images? It is the most widely used metric for evaluating generative models, and for good reason — a model that produces high-quality individual images but misses whole regions of the real distribution (mode collapse) will score poorly on FID even if individual samples look good.

To compute FID, both sets of images — real and generated — are passed through a pretrained Inception-v3 network (a deep CNN trained on ImageNet), and the activations from one of its intermediate layers are extracted. These activations form a high-dimensional feature space in which images with similar visual characteristics cluster together. We model each set's feature activations as a multivariate Gaussian distribution, characterized by a mean vector μ and covariance matrix Σ . The FID is then the Fréchet distance between these two Gaussians:

$$FID = \|\mu_r - \mu^{\wedge}\|^2 + \text{Tr}(\Sigma_r + \Sigma^{\wedge} - 2(\Sigma_r \Sigma^{\wedge})^{(1/2)})$$

Here, μ_r and Σ_r are the mean and covariance of the real image features, and $\mu^{\wedge}$ and $\Sigma^{\wedge}$ are those of the generated image features. $\text{Tr}(\cdot)$ is the matrix trace. The term $\|\mu_r - \mu^{\wedge}\|^2$ measures how far apart the centers of the two distributions are; the trace term measures how well their shapes match. A lower FID is better — a perfect generative model that exactly reproduces the real distribution would score zero. In practice, FID scores below 10 are considered excellent; scores above 50 suggest significant quality problems.

4.2 Structural Similarity Index Measure (SSIM)

While FID measures distributional similarity at the population level, the Structural Similarity Index Measure (SSIM) measures the similarity between two individual images. Introduced by Wang and colleagues [11], SSIM is motivated by a simple observation: the human visual system is particularly sensitive to structural information — the spatial arrangement of luminance patterns — rather than absolute pixel values. A metric that penalizes every slightly misaligned pixel equally (like mean squared error) misses this point.

SSIM decomposes image similarity into three components — luminance (l), contrast (c), and structure (s) — and combines them multiplicatively:

$$SSIM(x, y) = l(x, y) \cdot c(x, y) \cdot s(x, y)$$

where each component is computed over local image patches. The luminance term $l(x, y) = (2\mu_x\mu_y + C_1) / (\mu_x^2 + \mu_y^2 + C_1)$ compares mean pixel

values; the contrast term $c(x,y) = (2\sigma_x\sigma_y + C_2) / (\sigma_x^2 + \sigma_y^2 + C_2)$ compares standard deviations; and the structure term $s(x,y) = (\sigma_{xy} + C_3) / (\sigma_x\sigma_y + C_3)$ compares the normalized cross-correlation between patches. C_1, C_2, C_3 are small stabilization constants. The result is a value in $[-1, 1]$ where 1 means identical images, 0 means no structural correlation, and negative values indicate structural anti-correlation. For Chinese character generation, an SSIM above ~ 0.6 between a generated image and its nearest real-class neighbor indicates the generated character preserves the essential stroke structure.

4.3 Learned Perceptual Image Patch Similarity (LPIPS)

SSIM is a hand-crafted metric; Learned Perceptual Image Patch Similarity (LPIPS), introduced by Zhang and colleagues [12], asks instead: what does a deep neural network think about the similarity between two images? The idea is that a network trained on large amounts of image data has learned feature representations that correlate with human perceptual judgment far better than any hand-designed formula.

To compute LPIPS, both images are passed through a pretrained deep network (such as VGG — the Visual Geometry Group network — or AlexNet). The activations at multiple layers are extracted, normalized, and compared channel-by-channel. The differences are then aggregated into a single scalar distance:

$$LPIPS(x, y) = \sum_l w_l \|\varphi_l(x) - \varphi_l(y)\|^2$$

where $\varphi_l(\cdot)$ denotes the feature activations at layer l , and w_l are learned weights that calibrate the contribution of each layer. Unlike SSIM, LPIPS is a *distance* — lower is better, meaning the two images are perceptually more similar. A generated Chinese character with a low LPIPS distance to a real character of the same class not only has the right pixel structure, it looks right to a network that has learned something like visual perception.

4.4 Downstream CNN Classification Accuracy

All three of the above metrics evaluate the quality of generated images in isolation. But the ultimate test for our purposes is pragmatic: does adding the synthetic data actually help a classifier learn? We measure this by training a CNN classifier under two conditions — on real data alone (the baseline) and on real plus synthetic data (the augmented condition) — and comparing their validation accuracy. To ensure the comparison is statistically reliable, we repeat each training run five times with different random seeds and apply a paired t-test to determine whether any observed improvement is significant ($p < 0.05$) or within the noise of training variability. We also compute Cohen's d , a measure of effect size that tells us not just whether the difference is statistically significant, but how large it is in practical terms.

5. Method

5.1 Experimental Pipeline Overview

The experimental scaffold was designed to be architecture-agnostic, so that all three proposed architectures could be evaluated under identical conditions; the final study instantiates it for the C-WGAN-GP arm. (Two abbreviations appearing in the comparison table below — ELBO, the Evidence Lower Bound, and GP, the Gradient Penalty — are developed fully in Section 5.2.) The pipeline has four stages: (1) train the generative model on the real training set; (2) generate a synthetic dataset; (3) optionally filter the synthetic dataset through the quality pipeline; (4) train the CNN classifier on the combined real-plus-synthetic set and evaluate on held-out real data. Because everything outside the synthetic-data path is held constant, any observed difference in downstream accuracy is attributable to the synthetic data and its curation — in the executed experiments, to the presence and filtering of C-WGAN-GP samples — and not to any other experimental variable.

Figure 1 presents the experimental pipeline as a block diagram: real data is split once (stratified, seeded); the C-WGAN-GP trains on the real training subset; the trained Generator produces a large candidate pool which the trained Critic scores; the quality filter retains the top-scoring samples; and the CNN classifier is trained on real-plus-filtered-synthetic data and evaluated on the untouched real test set, with a real-only baseline path for comparison.

Figure 1 — Experimental pipeline: C-WGAN-GP generation, critic-based quality filtering, and downstream classification evaluation




Figure 1 — Experimental pipeline block diagram. Color coding: gold boxes are the datasets used for training and evaluation (real and augmented); green boxes are the generative components (C-WGAN-GP training, Generator, Critic); the red box is the two-stage quality filter, this paper’s original contribution; blue boxes are the evaluation components (CNN classifier and statistical testing); orange boxes are intermediate pipeline elements — the data split operation and the synthetic-data pools flowing into and out of the filter. Dashed arrows mark the critic-scoring path and the real-only baseline path.

Property	C-VAE	C-GAN	C-VAEGAN (Hybrid)
Training objective	ELBO (reconstruction + $\beta \cdot \text{KL divergence}$)	Minimax adversarial game	ELBO + adversarial perceptual loss
Latent space	Structured, continuous $N(0, I)$ prior	Unstructured noise vector z	Structured (VAE encoder) + GAN perceptual loss
Output quality	Smooth; may be slightly blurry	Sharp; risk of mode collapse	Sharp and structured (theoretical best)
Training stability	High — ELBO is well-defined	Medium — collapse risk without GP	Medium-High — GP regularization required
Key citations	[1], [2]	[3], [4], [7]	[5], [6]

5.2 Mathematical Foundations of Each Architecture

The C-VAE and the Evidence Lower Bound (ELBO)

The C-VAE rests on a beautiful probabilistic idea: rather than learning a fixed encoding for each image, learn a *distribution* over possible encodings. Formally, given an image x and its class label y , the encoder produces the parameters (μ, σ) of a Gaussian distribution $q\phi(z | x, y)$ over a latent vector z . The decoder then takes a sample z from this distribution (along with y) and reconstructs the image. The model is trained to maximize the Evidence Lower Bound (ELBO), which balances two competing objectives:

$$ELBO = E_{q\phi(z|x,y)} [\log p\theta(x | z, y)] - \beta \cdot KL(q\phi(z | x, y) || p(z))$$

The first term is the **reconstruction loss**: how well does the decoder reconstruct the original image from the sampled z ? The second term is the **KL (Kullback-Leibler) divergence**: how far is the learned encoding distribution $q\phi(z | x, y)$ from a standard normal distribution $p(z) = N(0, I)$? The KL term acts as a regularizer that keeps the latent space organized — preventing the model from using arbitrary, scattered encodings. The hyperparameter β controls the trade-off: higher β produces a smoother, more regular latent space at some cost to reconstruction sharpness.

The Reparameterization Trick

There is a subtle but critical engineering challenge lurking in the ELBO: to optimize the model with gradient descent, we need to backpropagate gradients through the *sampling step* $z \sim q\phi(z | x, y)$. But sampling is a stochastic operation, and stochastic operations are not differentiable in the usual sense. Gradients cannot flow through randomness.

The **reparameterization trick** is the elegant solution. Instead of sampling z directly from the learned distribution $N(\mu, \sigma^2)$, we rewrite the sampling operation as:

$$z = \mu + \sigma \odot \varepsilon, \varepsilon \sim N(0, I)$$

Here, ε is a sample from a *fixed* standard normal distribution — it carries all the randomness — while μ and σ are deterministic outputs of the encoder network. Think of it this way: imagine you need to draw a random point from a circle of radius r centered at point p . You could sample directly, but gradients would not flow. Instead, you sample a random direction ε from a unit circle, then *compute* the point as $p + r \cdot \varepsilon$. The randomness lives in ε , which we treat as a fixed input, not a parameter; all the learnable structure lives in p and r , through which gradients flow freely. The reparameterization trick shifts the randomness out of the computational graph and into a fixed noise term, making the entire operation differentiable with respect to μ and σ .

The C-GAN and Gradient Penalty

The C-GAN trains a Generator G and Discriminator D simultaneously with opposing objectives. Conditioned on class label y , G maps a noise vector z to a synthetic image $G(z | y)$; D takes an image (real or fake) and its label y and outputs a probability that the image is real. The minimax objective is:

$$\min_G \max_D [E[\log D(x | y)] + E[\log(1 - D(G(z | y)))]]$$

The Generator tries to minimize this (fool D); the Discriminator tries to maximize it (detect fakes). To stabilize training, we apply the Wasserstein Gradient Penalty (WGAN-GP) [7], which adds a term penalizing the Discriminator for having gradients too far from unit norm:

$$L^e += \lambda \cdot E[(\|\nabla^e D(\hat{x} | y)\|_2 - 1)^2]$$

where $\hat{x}$ is a random convex combination of real and generated images. This constraint keeps the Discriminator smooth, preventing the gradient signals sent to the Generator from exploding or vanishing — the root cause of most GAN training instability.

The C-VAEGAN Combined Objective

The C-VAEGAN uses the VAE encoder-decoder as its generative backbone, but replaces the VAE's pixel-level reconstruction loss with a perceptual loss computed from the GAN Discriminator's internal feature representations. The combined loss is:

$$L = L^{OMLN} + \lambda_a^{dv} \cdot L^{NAN} + \lambda^{feat} \cdot L^{feat,m}$$

The ELBO term maintains the structured latent space; the adversarial term sharpens outputs by asking the Discriminator whether they look real; the feature-matching term aligns the internal representations of real and reconstructed images within the Discriminator's layers, further improving perceptual quality.

5.3 Shared Architecture

To ensure a fair comparison, the design specifies a single convolutional backbone shared by all three proposed models; the C-WGAN-GP trained in the final experiments instantiates it. The encoder (used in C-VAE and C-VAEGAN) consists of four convolutional layers with Batch Normalization and LeakyReLU activations, reducing the input from 64×64 to a 4×4 feature map, then projecting to the latent parameters ($\mu, \log \sigma^2$). The class label is concatenated to the image as a tiled channel before the first convolutional layer. The decoder/generator uses four transposed convolutional layers (the “reverse” of convolution, used to upscale feature maps back to image size), producing a 64×64 grayscale output through a

Tanh activation that maps values to $[-1, 1]$. The discriminator uses the same four-layer convolutional structure as the encoder, ending in a scalar logit.

5.4 The Two-Stage Quality Filter (Original Contribution)

A critical and independently developed contribution of this paper is the two-stage quality filter, designed to curate synthetic samples before they enter the augmentation pool. The motivation is straightforward: generative models — particularly early in training or in underrepresented character classes — do not always produce good samples. Adding low-quality synthetic images to the training set could harm the classifier rather than help it. Filtering before augmentation ensures that only synthetic samples meeting minimum quality standards are used.

To the author’s knowledge, applying a two-stage discriminator-confidence-plus-perceptual-distance filter as a principled quality gate for generative data augmentation in the Chinese character domain has not been previously reported in the literature. The design was developed independently in the course of this project.

Stage 1 — Discriminator confidence: A synthetic image $\tilde{x}$ is accepted only if the trained Discriminator assigns it a confidence score above a threshold τ_D (approximately the top 70th percentile of generated samples). This stage filters out samples that the model itself “knows” are poor — images that look so far from the real distribution that even the Discriminator is not fooled.

Stage 2 — Perceptual proximity: Among samples passing Stage 1, we additionally require that the SSIM score between $\tilde{x}$ and its nearest real-class neighbor exceeds τ_S , and that the LPIPS distance is below τ_L . This stage filters out samples that may score well on the Discriminator but are perceptually dissimilar from any real example in the class — stylistically plausible but geometrically wrong. The thresholds are tuned on a held-out validation split.

The filter is designed to apply identically across architectures, enabling a clean A/B comparison between filtered and unfiltered augmentation. In the final study this comparison is realized for the C-GAN architecture, providing the test of H_3 reported in Section 6.

Implementation in the final study. Stage 1 was implemented and validated: for each class, a pool of 400 candidate images is generated and scored by the trained WGAN critic (conditioned on the class label), and the top 25% (100 images/class) are retained. Because the Wasserstein critic outputs an unbounded realism score rather than a probability, the threshold τ_D is realized as a per-class rank cutoff rather than an absolute confidence value — a natural adaptation that requires no calibration. Stage 2 (SSIM/LPIPS perceptual screening) was not implemented in the final

experiments and is carried forward as future work; the empirical results of Section 6 therefore constitute a validation of Stage 1 alone, and represent a lower bound on the filter’s potential benefit.

5.5 CNN Classifier

The evaluation classifier is a standard three-block CNN: each block consists of two Conv2d layers with Batch Normalization and ReLU activations followed by MaxPooling, reducing spatial dimensions from 64×64 down to an 8×8 feature map. Two fully connected layers with Dropout ($p = 0.4$) and a softmax output over 15 classes complete the architecture. The same classifier is used in all conditions — the baselines and every augmented variant (unfiltered and critic-filtered, at 1:1 and 4:1 synthetic-to-real ratios) — trained with the Adam optimizer and a cosine learning rate schedule. Holding the classifier fixed ensures that differences in accuracy between conditions reflect the quality of the augmentation, not the classifier design.

5.6 Hyperparameters Used

The midpoint proposal specified a hyperparameter search across latent dimension, KL weight, gradient-penalty coefficient, and augmentation ratio. Under the descoped final design, fixed literature-standard values were used for the generative model, and the augmentation-ratio dimension of the search was retained (1:1 and 4:1 synthetic-to-real evaluated in both experiments). The fixed values follow the WGAN-GP reference settings of Gulrajani et al. [7]:

Parameter	Value	Rationale
Latent dimension d	128	Upper midpoint candidate; capacity for 15-class conditional generation
Gradient penalty λ	10	WGAN-GP reference default [7]
Critic steps per generator step	5	WGAN-GP reference default [7]
Optimizer (G and C)	Adam, lr $1e-4$, $\beta_1 = 0.0$, $\beta_2 = 0.9$	WGAN-GP reference settings [7]
WGAN-GP training epochs	500	Selected after a 100-epoch pilot showed non-converged generator loss
CNN: epochs / lr / dropout	30 / $1e-3$ (cosine) / 0.4	Fixed across all conditions for comparison fairness

Parameter	Value	Rationale
Augmentation ratio (searched)	1:1, 4:1	Evaluated in both experiments
Filter pool / retention	400/class, top 25%	Stage 1 quality filter operating point (not tuned)

A full search over the generative hyperparameters and the filter retention threshold is deferred to future work; Section 7 discusses the threshold sweep as a priority follow-up.

6. Experimental Results

Two experiments were conducted, sharing the identical stratified split, five random seeds (42-46), CNN classifier, and held-out real test set (100 images/class, 1,500 total). All significance tests are paired t-tests across seeds against the within-experiment baseline, $\alpha = 0.05$.

Experiment 1 — unfiltered augmentation under mild scarcity. The C-WGAN-GP was trained for 500 epochs on 50 real images/class (750 total). Synthetic images were sampled without curation and mixed with the real training data at 1:1 and 4:1 synthetic-to-real ratios.

Experiment 2 — critic-filtered augmentation under deep scarcity. The real training set was reduced to 25 images/class (a strict subset of Experiment 1’s training set, leaving val/test untouched). Synthetic data was drawn from the same 500-epoch generator checkpoint, but passed through the Stage 1 quality filter: 400 candidates/class scored by the trained critic, top 25% retained.

6.1 Headline Results

Experiment	Condition	Mean Acc	Std	Δ vs Baseline	p-value	Significant?
Exp 1 (50/class, unfiltered)	Baseline	88.49%	0.78%	—	—	—
	Augmented 1:1	87.44%	0.89%	−1.05%	0.057	No
	Augmented 4:1	86.27%	0.25%	−2.22%	0.008	Yes (−)
Exp 2 (25/class)	Baseline	71.77%	2.57%	—	—	—

Experiment	Condition	Mean Acc	Std	Δ vs Baseline	p-value	Significant?
	filtered)					
	Augmented 1:1	77.32%	1.10%	+5.55%	0.0085	Yes (+)
	Augmented 4:1	80.19%	1.63%	+8.42%	0.0005	Yes (+)

Figure 2 — Classification accuracy across all six conditions (mean $\pm$ std, 5 seeds). Unfiltered augmentation degrades accuracy; critic-filtered augmentation improves it.




Figure 2 — Classification accuracy across all six conditions.

The direction of the augmentation effect reverses between experiments. Unfiltered synthetic data degrades accuracy, and the degradation grows with the number of synthetic samples. Critic-filtered synthetic data improves accuracy, and the improvement grows with the number of synthetic samples — every individual seed improved in both filtered conditions. The filtered 4:1 condition recovers nearly half the accuracy lost to the halving of real training data (71.77% $\rightarrow$ 80.19%, against 88.49% with 50 real images/class).

6.2 Training Dynamics

The C-WGAN-GP exhibited the loss behavior characteristic of a healthy gradient-penalty run: large early oscillations (epochs 1–20), followed by stabilization, with the critic loss settling near -4.8 and the generator loss plateauing in the 6.0–6.8 band from roughly epoch 200 onward. A 100-epoch pilot run, by contrast, ended with the generator loss still rising — and synthetic samples from that checkpoint, when used for augmentation, produced a 10.96-percentage-point accuracy *drop* (77.05% vs the 88.01% baseline measured in the pilot configuration). Extending training to 500

epochs reduced unfiltered augmentation’s harm to the -1% to -2% range reported above. Generator convergence is thus necessary but, in this study, not sufficient: only the addition of the quality filter converted augmentation from harmful to beneficial.

Classifier training loss provides the mechanism’s signature. In the unfiltered 4:1 condition, final-epoch training loss fell to 0.01–0.02 (versus 0.09–0.28 for the baseline) with the lowest seed variance of any condition — the classifier reliably memorized the synthetic distribution’s artifacts rather than learning transferable character structure. In the filtered conditions, training loss remained in a healthy intermediate range and test variance *fell* relative to the scarce baseline (1.10% vs 2.57% at 1:1), indicating the filtered synthetic data acted as a stabilizer as well as a signal source.

6.3 Quality Filter Analysis

Visual inspection across the filter boundary confirms the critic ranks usefully: retained samples show cleaner stroke structure and fewer brightness artifacts than the unfiltered population, while rejected samples include the spatially discontinuous and low-contrast failures characteristic of the generator’s weaker outputs. The filter’s operating point (top 25% of a 400/class pool) was chosen a priori and not tuned; the relationship between retention threshold and downstream benefit is unmeasured and is the highest-priority follow-up experiment after the architecture comparison.

6.4 Status of the Designed Metric Suite

The generation-quality metrics specified in Section 4 (FID, SSIM, LPIPS) were not computed in the final experiments; downstream classification accuracy — the metric this paper argues is the decision-relevant one for the augmentation use case — served as the sole quantitative criterion. Computing the full metric suite for retained versus rejected filter samples would link the filter’s behavior to measurable image statistics and is listed in future work.

7. Conclusions

This study set out to determine whether conditional generative models can usefully augment scarce training data for Chinese character classification, with the question sharpened at midpoint into three hypotheses. The verdicts:

H₁ (Generation Quality, cross-architecture): *Not tested.* The three-architecture comparison was descoped; H₁ requires the C-VAE and C-VAEGAN implementations and the FID/SSIM/LPIPS metric suite, all carried forward as future work.

H₂ (Augmentation Efficacy): *Conditionally supported — and the condition is the finding.* For the C-WGAN-GP, unfiltered synthetic augmentation failed H₂, significantly degrading accuracy at a 4:1 ratio (−2.22%, $p = 0.008$). Critic-filtered augmentation under deeper scarcity strongly satisfied H₂’s significance criterion (+8.42% at 4:1, $p = 0.0005$), with accuracy improving monotonically with the fraction of augmentation samples.

H₃ (Quality Filter Value): *Supported for the C-GAN architecture.* The Stage 1 filter — repurposing the trained critic as a quality oracle over a generated candidate pool — is precisely what separated harmful from beneficial augmentation in this study. The cross-architecture portion of H₃ awaits the comparative study.

An important caveat qualifies the headline reversal: the two experiments differ in both scarcity level and filtering, so the sign flip cannot be attributed to the filter alone from this design — the within-experiment ratio trends are each internally clean, but the 2×2 factorial (filtering × scarcity) needed for clean attribution was not run. It is plausible that part of the filtered conditions’ benefit derives from the greater headroom at 25 images/class.

The practical lesson for low-resource domains — including the medical imaging setting this study proxies — is two-sided. Generative augmentation deployed naively can measurably harm the downstream model, even from a converged, visually plausible generator. But the remedy is nearly free: the adversarial critic, normally discarded after training, ranks the generator’s own outputs well enough that simply keeping the top quartile of a candidate pool converted a −2.22% effect into a +8.42% effect in this study. *Train the pair, keep the critic, score the candidates, train on the best.*

Future work, in priority order: (1) the 2×2 factorial to attribute the reversal; (2) a retention-threshold sweep to map the quality-quantity tradeoff; (3) Stage 2 of the filter (SSIM/LPIPS perceptual screening); (4) the original C-VAE / C-GAN / C-VAEGAN comparison under the identical filtered protocol, testing H₁ and completing H₃; (5) the FID/SSIM/LPIPS metric suite computed for retained versus rejected samples; (6) transfer to a genuine medical imaging task and to a curated CASIA-HWDB subset.

8. Computational Resources Used

All experiments ran on a single NVIDIA L4 Graphics Processing Unit (GPU) via Google Colab Pro. The dominant cost was the 500-epoch C-WGAN-GP training run (≈3.5 hours); the 100-epoch pilot, thirty CNN classifier training runs (six conditions × five seeds, 2–4 minutes each), and synthetic generation/filtering brought the total to approximately 8–10 GPU-hours — well under the 120 GPU-hours estimated at midpoint for the full three-architecture study, consistent with the descoped design. Model checkpoints

were saved to persistent storage every 25 epochs, enabling recovery across session interruptions and exact reuse of the trained generator and critic across both experiments. All code is PyTorch; the complete pipeline, with per-cell documentation and all executed outputs, accompanies this paper as a Jupyter notebook.

Acknowledgements

This paper emerged from several months of sustained, iterative dialogue with Claude (Anthropic), used here as a research thinking partner rather than a writing tool. The core research question, the choice of Chinese characters as a domain, the three-model comparative framing, and the experimental hypotheses all originated in my own thinking; Claude's role was to sharpen those ideas through rigorous back-and-forth, flag gaps in reasoning, suggest implementation strategies, and help translate rough intuitions into precise technical language. The architecture decisions, the quality filter pipeline, and the dataset selection rationale each went through multiple rounds of challenge and refinement in that dialogue. I find this mode of working — bringing your own ideas and using AI to stress-test and develop them — to be a genuinely productive form of intellectual partnership, and wanted to acknowledge it honestly rather than obscure it.

References

- [1] D. P. Kingma and M. Welling, "Auto-Encoding Variational Bayes," in Proc. ICLR, 2014. arXiv:1312.6114.
- [2] K. Sohn, H. Lee, and X. Yan, "Learning Structured Output Representation using Deep Conditional Generative Models," in Proc. NeurIPS, vol. 28, 2015.
- [3] I. J. Goodfellow et al., "Generative Adversarial Nets," in Proc. NeurIPS, vol. 27, 2014. arXiv:1406.2661.
- [4] M. Mirza and S. Osindero, "Conditional Generative Adversarial Nets," arXiv:1411.1784, 2014.
- [5] A. B. L. Larsen, S. K. Sønderby, H. Larochelle, and O. Winther, "Autoencoding beyond Pixels using a Learned Similarity Metric," in Proc. ICML, 2016. arXiv:1512.09300.
- [6] B. Kong and Y. Xu, "Generative Adversarial Networks for Chinese Character Image Synthesis," unpublished manuscript, 2021.
- [7] I. Gulrajani, F. Ahmed, M. Arjovsky, V. Dumoulin, and A. Courville, "Improved Training of Wasserstein GANs," in Proc. NeurIPS, vol. 30, 2017. arXiv:1704.00028.

- [8] C. Zhang, P. Yin et al., "Chinese Character Recognition with Deep CNNs," in Proc. ICDAR, 2017.
- [9] C.-L. Liu, F. Yin, D.-H. Wang, and Q.-F. Wang, "CASIA Online and Offline Chinese Handwriting Databases," in Proc. ICDAR, 2011.
- [10] Chinese-MNIST Dataset. Kaggle. [Online]. Available: <https://www.kaggle.com/datasets/gpreda/chinese-mnist>.
- [11] Z. Wang, A. C. Bovik, H. R. Sheikh, and E. P. Simoncelli, "Image Quality Assessment: From Error Visibility to Structural Similarity," IEEE Trans. Image Process., vol. 13, no. 4, pp. 600–612, 2004.
- [12] R. Zhang et al., "The Unreasonable Effectiveness of Deep Features as a Perceptual Metric," in Proc. CVPR, 2018. arXiv:1801.03924.
- [13] M. Arjovsky, S. Chintala, and L. Bottou, "Wasserstein GAN," in Proc. ICML, 2017. arXiv:1701.07875.